



Deep Learning: Autoencoder and Restricted Boltzmann Machine

Bo Peng

Definition

- Deep learning means using a neural network with several layers of nodes between inputs and outputs
- The series of layers between input & output do feature identification and processing in a series of stages, just as our brains seem to.
- Shallow vs Deep
 - Representation of functions of inputs
 - Compact representation: a function can be compactly represented by a deep architecture, it might need a very large architecture to be represented by an insufficient deep one. (Bengio, 2009)

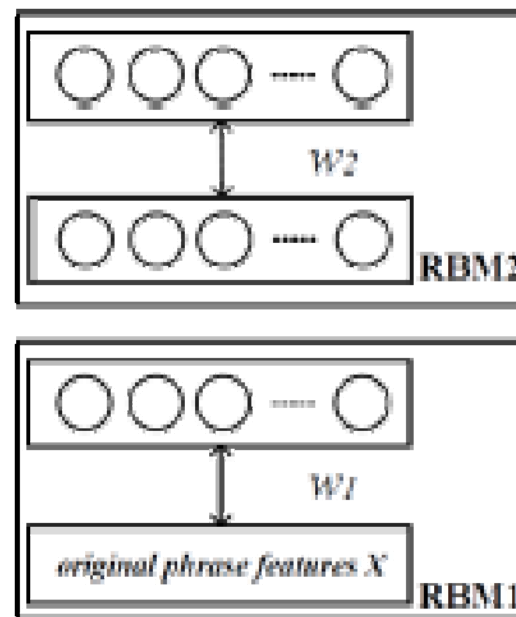


Breakthroughs


- Problem of multiple-layer neural network
 - Hard to train with too many parameters
 - SGD highly relies on “good” initial weights
 - Large initial weights: pool local minima
 - Small initial weights: gradients are tiny
 - Unsupervised learning?
- Breakthroughs
 - Deep Belief Network (DBN)
Hinton, G.E, Osindero, S., and Teh, Y.W. (2006).
A fast learning algorithm for deep belief nets.
 - Autoencoders
Bengio, Y., Lamblin, P., Popovici, P., Larochelle, H. (2007)
Greedy Layer-Wise Training of Deep Networks.

Intuitive Idea

- What did Hinton and Bengio do?
 - Model a two-layer network at a time
 - Pretraining procedure for initial weights
 - Global fine tuning stage
 - Backpropagation



(a) Pre-training

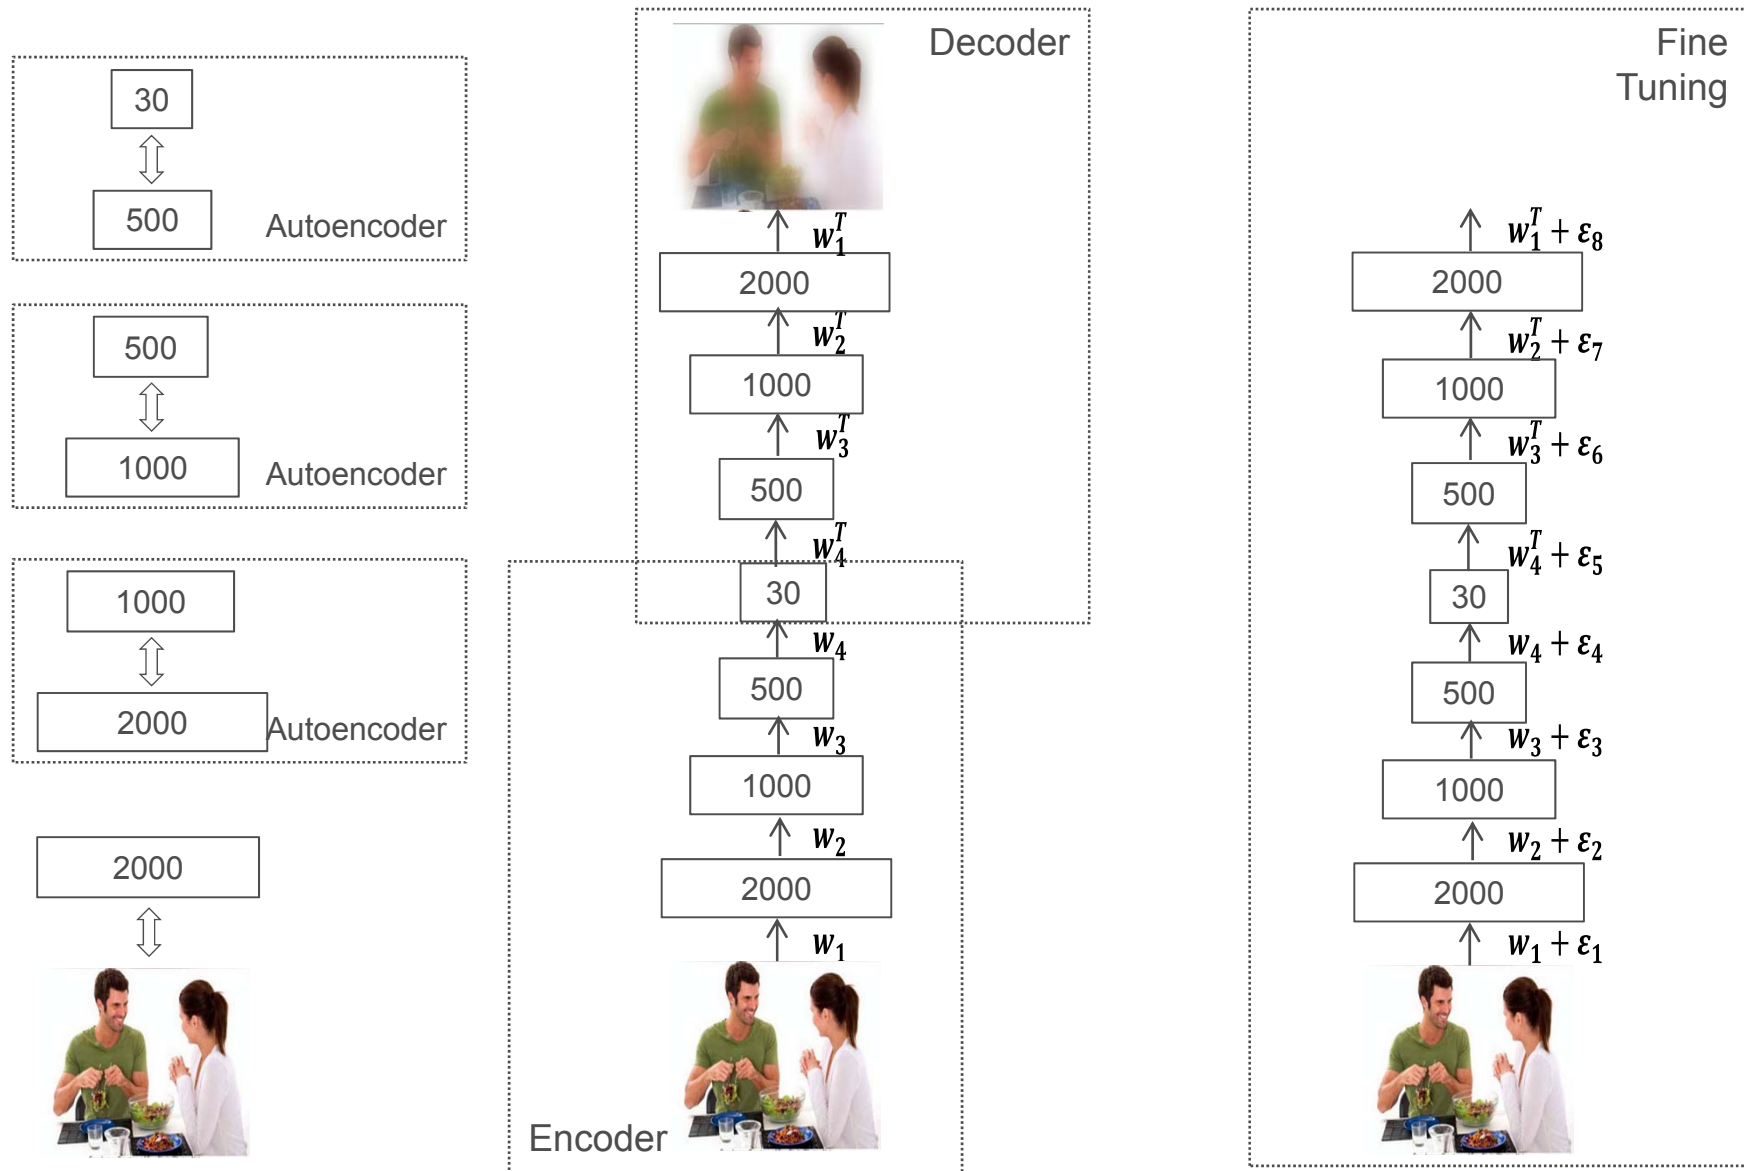


(b) DBN

Autoencoder

- What?
 - Deterministic artificial neural network
 - Learn a compressed distributed representation (encoding) for a set of data
 - Dimension reduction
- How?
 - Feed-forward pass to compute activation at all hidden layers
Encode: $y = s(Wx + b)$ where s is sigmoid or hyperbolic tangent
 - Measure the deviation of output from the input (MSE)
Decode: $z = s(W'y + b')$ with $W' = W^T$
Reconstruction error: $L(xz) = ||x - z||^2$ or $L(xz) = -\sum_{k=1}^d [x_k \log z_k + (1 - x_k) \log(1 - z_k)]$
 - Backpropagate the error through the net for weights updating

Stacked Autoencoder



Some issues

- Pretraining
 - Unsupervised initialization in a greedy layer-wise fashion will put the parameters in a region of parameters space from which a good local optimum can be achieved by local descent
- Fine tune
 - Supervised learning: Classification
 - Unsupervised learning: Dimension reduction
- Denoising
 - Robust feature extraction and avoid simple copy
 - Input corruption

Restricted Boltzmann Machine

- What?
 - Stochastic artificial neural network
 - Learn a probability distribution over a set of inputs
 - Dimension reduction and classification

- Energy-Based Models (EBM)


- $p(x) = \frac{e^{-E(x)}}{Z}$ where $Z = \sum_x e^{-E(x)}$
 - Hidden Units

$$P(x) = \sum_h P(x, h) = \sum_h \frac{e^{-E(x)}}{Z}$$

Set *free energy* $F(x) = -\log \sum_h e^{-E(x)}$, we have $P(x) = \frac{e^{-E(x)}}{Z}$

- Gradient
- $$-\frac{\partial \log p(x)}{\partial \theta} = \frac{\partial F(x)}{\partial \theta} - \sum_{\tilde{x}} p(\tilde{x}) \frac{\partial F(\tilde{x})}{\partial \theta}$$

Restricted Boltzmann Machine



- log-linear Markov Random Field (MRF)
 - Energy function is linear in its free parameters

$$E(v, h) = -b'v - c'h - h'Wv$$

$$F(v) = -b'v - \sum_i \log \sum_{h_i} e^{h_i(c_i + W_i v)}$$

- No visible-visible and hidden-hidden connections

$$p(h|v) = \prod_i p(h_i|v)$$

$$p(v|h) = \prod_j p(v_j|h)$$



Restricted Boltzmann Machine

- Binary Units: v_j and $h_i \in \{0,1\}$

$$F(v) = -b'v - \sum_i \log(1 + e^{(c_i + W_i v)})$$

$$P(h_i = 1|v) = \text{sigm}(c_i + W_i v)$$

$$P(v_j = 1|h) = \text{sigm}(b_j + W_j' h)$$

- Updates

$$-\frac{\partial \log p(v)}{\partial W_{ij}} = E_v[p(h_i|v) \cdot v_j] - v_j^{(i)} \cdot \text{sigm}(W_i \cdot v^{(i)} + c_j)$$

$$-\frac{\partial \log p(v)}{\partial c_i} = E_v[p(h_i|v)] - \text{sigm}(W_i \cdot v^{(i)})$$


$$-\frac{\partial \log p(v)}{\partial b_j} = E_v[p(v_j|h)] - v_j^{(i)}$$

Restricted Boltzmann Machine

- Gibbs sampling

$$h^{(n+1)} \sim \text{sigm}(W'v^{(n)} + c)$$

$$v^{(n+1)} \sim \text{sigm}(Wh^{(n+1)} + b)$$




Deep Belief Networks

- Stacked RBM


$$P(x, h^1, \dots, h^l) = \left(\prod_{k=0}^{l-2} P(h^k | h^{k+1}) \right) P(h^{l-1}, h^l)$$

where $x = h^0$, $P(h^{k-1} | h^k)$ is a conditional distribution for visible units conditioned on the hidden units of the RBM at level k , and $P(h^{l-1}, h^l)$ is the visible-hidden joint distribution in the top-level RBM.



Example


- Data description
 - MNIST database of handwritten digits
 - $28 \times 28 = 784$ pixels for each image
 - Training set 60,000 images and testing set 10,000 images
- Python
 - Deep learning library (Bengio): Theano
 - dA, sda, rbm, DBN (<http://deeplearning.net/tutorial/>)




Results(More nodes)

Autoencoder


Layer: 30 MSE: 0.05074




Layer: 100 MSE: 0.02686



Layer: 500 MSE: 0.00786



Layer: 1000 MSE: 0.00941




RBM


Layer: 30 MSE: 0.06213




Layer: 100 MSE: 0.02350



Layer: 500 MSE: 0.00977




Layer: 1000 MSE: 0.00973



Results(More layers)

Autoencoder

Layer: 30 MSE: 0.05074




Layer: 250-30 MSE: 0.05450



Layer: 500-250-30 MSE: 0.04632



Layer: 1000-500-250-30 MSE: 0.05012




RBM


Layer: 30 MSE: 0.06213




Layer: 250-30 MSE: 0.04957



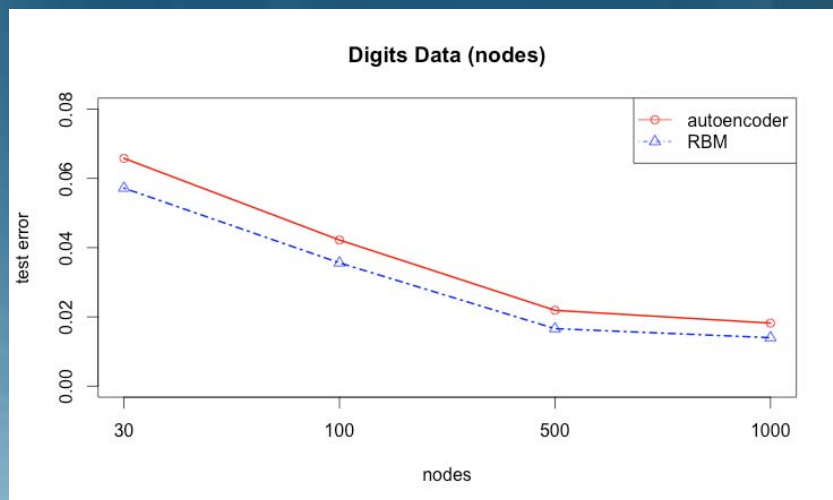
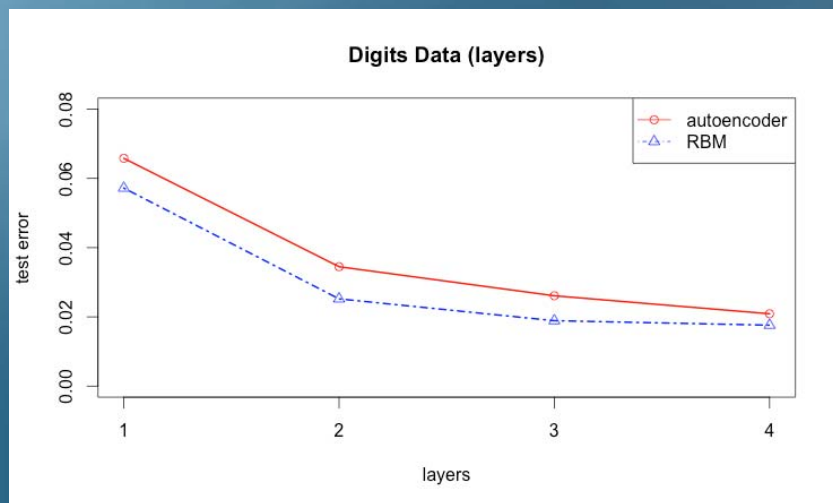
Layer: 500-250-30 MSE: 0.03741



Layer: 1000-500-250-30 MSE: 0.0476



Results(Classification)



- RBM performs better on classification than autoencoder
- Test error decreases as layers and nodes size increase

Reference

- Reducing the dimensionality of data with Neural networks. (Hinton and Salakhutdinov, 2006)
- A fast learning algorithm for deep belief nets. (Hinton, Osindero and Teh, 2006)
- Learning deep architectures for AI. (Bengio, 2009)